

Bonus Question:

Experiment with the hyperparameter patience (you may set it equal to 2 and 4 for instance) in the `fit_with_early_stopping()` function, and report your observations

Answer:

The patience affects how much allowance the program gives for an increase in validation loss. For example, a patience of 1 stopped the program immediately after just 1 increase in validation loss. This resulted in a very early stop after just 2 epochs. The validation accuracy wasn't as good as when patience was > 1 . **Patience = 2 produced the best results** with a validation accuracy of 0.7650390863418579 with the second-best accuracy being 0.764453113079071 with patience = 3.

Patience = 1:

```
{'val_loss': 1.0575330257415771, 'val_acc': 0.7626953125}
```

```
Epoch [0], train_loss: 0.2081, val_loss: 0.9186, val_acc: 0.7729
```

```
Epoch [1], train_loss: 0.1446, val_loss: 1.0705, val_acc: 0.7663
```

Early stopping at epoch 1

Patience = 2:

```
{'val_loss': 0.9376282691955566, 'val_acc': 0.7650390863418579}
```

```
Epoch [0], train_loss: 1.7795, val_loss: 1.4387, val_acc: 0.4651
```

```
Epoch [1], train_loss: 1.2620, val_loss: 1.1309, val_acc: 0.5972
```

```
Epoch [2], train_loss: 1.0170, val_loss: 0.9628, val_acc: 0.6573
```

```
Epoch [3], train_loss: 0.8562, val_loss: 0.8607, val_acc: 0.6984
```

```
Epoch [4], train_loss: 0.7315, val_loss: 0.8193, val_acc: 0.7169
```

```
Epoch [5], train_loss: 0.6205, val_loss: 0.7509, val_acc: 0.7534
```

```
Epoch [6], train_loss: 0.5309, val_loss: 0.7286, val_acc: 0.7582
```

```
Epoch [7], train_loss: 0.4505, val_loss: 0.7181, val_acc: 0.7620
```


```
Epoch [8], train_loss: 0.3645, val_loss: 0.7590, val_acc: 0.7661
```

```
Epoch [9], train_loss: 0.2936, val_loss: 0.7482, val_acc: 0.7766
```

```
Epoch [10], train_loss: 0.2453, val_loss: 0.8385, val_acc: 0.7710
```

```
Epoch [11], train_loss: 0.1871, val_loss: 0.9264, val_acc: 0.7751
```

Early stopping at epoch 11



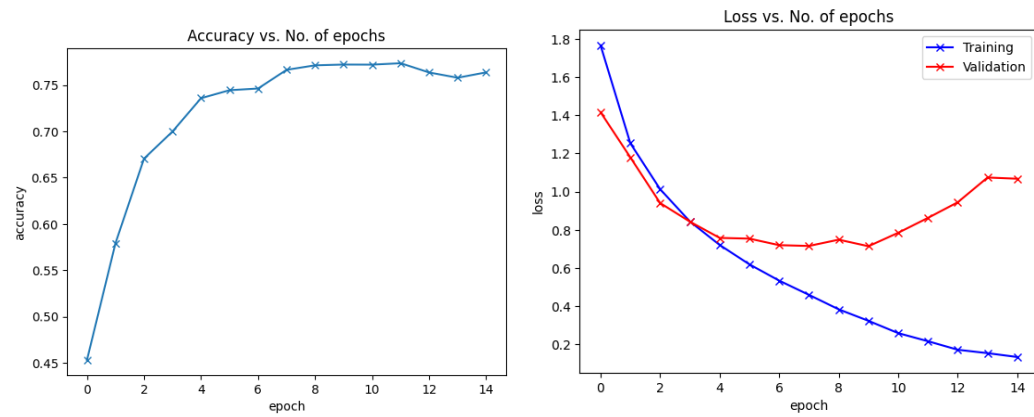
Patience = 3:

```
{'val_loss': 1.0728938579559326, 'val_acc': 0.764453113079071}
```

```
history_with_early_stopping = fit_with_early_stopping(num_epochs, lr, model3, train_dl, val_dl,
opt_func)
```

output

```
Epoch [0], train_loss: 1.7655, val_loss: 1.4156, val_acc: 0.4529
Epoch [1], train_loss: 1.2532, val_loss: 1.1793, val_acc: 0.5792
Epoch [2], train_loss: 1.0130, val_loss: 0.9405, val_acc: 0.6706
Epoch [3], train_loss: 0.8440, val_loss: 0.8431, val_acc: 0.6999
Epoch [4], train_loss: 0.7211, val_loss: 0.7579, val_acc: 0.7359
Epoch [5], train_loss: 0.6195, val_loss: 0.7542, val_acc: 0.7445
Epoch [6], train_loss: 0.5343, val_loss: 0.7200, val_acc: 0.7463
Epoch [7], train_loss: 0.4599, val_loss: 0.7158, val_acc: 0.7665
Epoch [8], train_loss: 0.3841, val_loss: 0.7487, val_acc: 0.7714
Epoch [9], train_loss: 0.3239, val_loss: 0.7143, val_acc: 0.7722
Epoch [10], train_loss: 0.2595, val_loss: 0.7848, val_acc: 0.7721
Epoch [11], train_loss: 0.2172, val_loss: 0.8628, val_acc: 0.7737
Epoch [12], train_loss: 0.1725, val_loss: 0.9444, val_acc: 0.7636
Epoch [13], train_loss: 0.1548, val_loss: 1.0740, val_acc: 0.7580
Epoch [14], train_loss: 0.1346, val_loss: 1.0676, val_acc: 0.7638
Early stopping at epoch 14
```



```
[75] # -----
# Graded Cell
# -----
dog_dir = "/train/dog"
dog_files = os.listdir(data_dir + dog_dir)
len_dog = len(dog_files)
print('No. of training examples for dogs', len_dog)
print(dog_files[:5])
```

```
# -----
# Graded Cell
# -----
cat_test_dir = "/test/cat"
cat_test_files = os.listdir(data_dir + cat_test_dir)
len_test_cat = len(cat_test_files)
print("No. of test examples for cats:", len_test_cat)
print(cat_test_files[:5])

No. of test examples for cats: 1000
['0414.png', '0646.png', '0328.png', '0129.png', '0628.png']
```

✓
0s

```
[80] # -----  
# Graded Cell  
# -----  
  
def length(training_data): # outputs the no. of images contained in the training set  
    # one line of code  
    return len(training_data)  
  
length(training_data)  
  
50000
```

✓
0s




```
# -----  
# Graded Cell  
# -----  
  
def shape(img): # outputs the shape of an image tensor  
    # one line of code  
    return img.shape  
  
img, label = training_data[0]  
img_size = shape(img)  
print(img_size, label)  
img  
  
torch.Size([3, 32, 32]) 0  
tensor([[0.7922, 0.7922, 0.8000, ..., 0.8118, 0.8039, 0.7961],  
        [0.8078, 0.8078, 0.8118, ..., 0.8235, 0.8157, 0.8078],  
        [0.8235, 0.8275, 0.8314, ..., 0.8392, 0.8314, 0.8235],  
        ...,  
        [0.8549, 0.8235, 0.7608, ..., 0.9529, 0.9569, 0.9529],  
        [0.8588, 0.8510, 0.8471, ..., 0.9451, 0.9451, 0.9451],  
        [0.8510, 0.8471, 0.8510, ..., 0.9373, 0.9373, 0.9412]],  
        [[0.8000, 0.8000, 0.8078, ..., 0.8157, 0.8078, 0.8000],  
         [0.8157, 0.8157, 0.8196, ..., 0.8275, 0.8196, 0.8118],  
         [0.8314, 0.8353, 0.8392, ..., 0.8392, 0.8353, 0.8275],  
         ...,  
         [0.8510, 0.8196, 0.7608, ..., 0.9490, 0.9490, 0.9529],  
         [0.8549, 0.8471, 0.8471, ..., 0.9412, 0.9412, 0.9412],  
         [0.8471, 0.8431, 0.8471, ..., 0.9333, 0.9333, 0.9333]],  
        [[0.7804, 0.7804, 0.7882, ..., 0.7843, 0.7804, 0.7765],  
         [0.7961, 0.7961, 0.8000, ..., 0.8039, 0.7961, 0.7882],  
         [0.8118, 0.8157, 0.8235, ..., 0.8235, 0.8157, 0.8078],  
         ...,  
         [0.8706, 0.8392, 0.7765, ..., 0.9686, 0.9686, 0.9686],  
         [0.8745, 0.8667, 0.8627, ..., 0.9608, 0.9608, 0.9608],  
         [0.8667, 0.8627, 0.8667, ..., 0.9529, 0.9529, 0.9529]]])
```

✓
0s

```
# -----  
# Graded Cell  
# -----  
show_example(*training_data[1000])
```


Label: airplane (0)



✓
0s

```
[87] # -----  
# Graded Cell  
# -----  
show_example(*training_data[2000])
```

Label: airplane (0)



✓
Os

```
[89] # -----  
# Graded Cell  
# -----  
def get_train_size(training_data, val_size): # outputs the size of the validation set  
    train_length = length(training_data) - val_size  
    return train_length  
  
val_size = 5000  
train_size = get_train_size(training_data, val_size)  
  
train_ds, val_ds = random_split(training_data, [train_size, val_size])  
len(train_ds), len(val_ds)  
  
(45000, 5000)
```

✓
Os

```
[94] # -----  
# Graded Cell  
# -----  
def apply_filter(image, filter):  
    img_rows, img_cols = image.shape # image dimensions  
    fil_rows, fil_cols = filter.shape # filter dimensions  
    out_rows, out_cols = (img_rows - fil_rows + 1, img_cols - fil_cols + 1) # output dimensions  
    output = torch.zeros([out_rows, out_cols])  
    for i in range(out_rows):  
        for j in range(out_cols):  
            output[i,j] = torch.sum(image[i:i+fil_rows,j:j+fil_cols] * filter)  
    return output
```

✓
0s



```
# -----  
# Graded Cell  
# -----  
class ImageClassificationBase(nn.Module):  
    def training_step(self, batch):  
        images, labels = batch  
        #####  
        # forward pass  
        out = self(images) # Generate predictions for all the images in the batch  
        loss = F.cross_entropy(out, labels) # Calculate cross-entropy (CE) loss  
        #####  
        return loss  
  
    def validation_step(self, batch):  
        images, labels = batch  
        #####  
        # validation pass  
        out = self(images) # Generate predictions for all the images in the batch  
        loss = F.cross_entropy(out, labels) # Calculate cross-entropy (CE) loss  
        #####  
        acc = accuracy(out, labels) # Calculate accuracy  
        return {'val_loss': loss.detach().cpu(), 'val_acc': acc}  
  
    def validation_epoch_end(self, outputs):  
        batch_losses = [x['val_loss'] for x in outputs] # Creates a list that contains the validation  
        epoch_loss = torch.stack(batch_losses).mean() # Combines batch losses into a single tensor  
        batch_accs = [x['val_acc'] for x in outputs] # Creates a list that contains the validation  
        epoch_acc = torch.stack(batch_accs).mean() # Combines batch accuracies into a single tensor  
        return {'val_loss': epoch_loss.item(), 'val_acc': epoch_acc.item()}  
  
    def epoch_end(self, epoch, result):  
        print("Epoch [{}], train_loss: {:.4f}, val_loss: {:.4f}, val_acc: {:.4f}".format(  
            epoch, result['train_loss'], result['val_loss'], result['val_acc']))  
  
    def accuracy(outputs, labels):  
        #####  
        # generate predictions (see HINT below)  
        _, preds = torch.max(outputs, dim=1) # Wrote this  
        # HINT: Use torch.max(input, dim) specifying input and dim while having in mind that torch.max() r  
        # along the specified dimension and the indices of those maximum values. Here, we are interested i  
        #####  
        return torch.tensor(torch.sum(preds == labels).item() / len(preds)).detach().cpu()
```

✓
0s

```
# -----  
# Graded Cell  
# -----  
class Cifar10CnnModel(ImageClassificationBase):  
    def __init__(self):  
        super().__init__()  
        self.network = nn.Sequential(  
            # input: 3 x 32 x 32  
            nn.Conv2d(3, 32, kernel_size=3, stride=1, padding=1),  
            # output: 32 x 32 x 32  
            nn.ReLU(),  
            # output: 32 x 32 x 32  
            nn.Conv2d(32, 64, kernel_size=3, stride=1, padding=1),  
            # output: 64 x 32 x 32  
            nn.ReLU(),  
            # output: 64 x 32 x 32  
            nn.MaxPool2d(2, 2),  
            # output: 64 x 16 x 16  
  
            nn.Conv2d(64, 128, kernel_size=3, stride=1, padding=1),  
            # output: 128 x 16 x 16  
            nn.ReLU(),  
            # output: 128 x 16 x 16  
            nn.Conv2d(128, 128, kernel_size=3, stride=1, padding=1),  
            # output: 128 x 16 x 16  
            nn.ReLU(),  
            # output: 128 x 16 x 16  
            nn.MaxPool2d(2, 2),  
            # output: 128 x 8 x 8  
  
            nn.Conv2d(128, 256, kernel_size=3, stride=1, padding=1),  
            # output: 256 x 8 x 8  
            nn.ReLU(),  
            # output: 256 x 8 x 8  
            nn.Conv2d(256, 256, kernel_size=3, stride=1, padding=1),  
            # output: 256 x 8 x 8  
            nn.ReLU(),  
            # output: 256 x 8 x 8  
            nn.MaxPool2d(2, 2),  
            # output: 256 x 4 x 4  
  
            nn.Flatten(),  
            nn.Linear(4096, 1024),  
            # output: 1024  
            nn.ReLU(),  
            # output: 1024  
            nn.Linear(1024, 512),  
            # output: 512  
            nn.ReLU(),  
            # output: 512  
            nn.Linear(512, 10))  
            # output: 10  
  
    def forward(self, xb):  
        return self.network(xb)
```



```
✓ 0s # -----
# Graded Cell
# -----
@torch.no_grad()
def evaluate(model, val_loader):
    model.eval()
    outputs = [model.validation_step(batch) for batch in val_loader]
    return model.validation_epoch_end(outputs)

def fit(epochs, lr, model, train_loader, val_loader, opt_func=torch.optim.SGD): # SGD refers to mini-b
    history = []
    optimizer = opt_func(model.parameters(), lr)
    for epoch in range(epochs):
        # Training Phase
        model.train()
        train_losses = []
        for batch in train_loader:
            loss = model.training_step(batch)
            #####
            # Arrange the following three lines in the correct order:
            # 1. optimizer.step(); 2. optimizer.zero_grad(); 3. loss.backward().
            loss.backward() # Calculate losses
            optimizer.step() # Step
            optimizer.zero_grad() # Reset gradients
            #####
            train_losses.append(loss.detach().cpu())
        # Validation phase
        result = evaluate(model, val_loader)
        result['train_loss'] = torch.stack(train_losses).mean().item()
        model.epoch_end(epoch, result)
        history.append(result) # displays train_loss, val_loss and val_acc at the end of each epoch
    return history
```

```
✓ 5s [125] test_loader = DeviceDataLoader(DataLoader(test_dataset, batch_size*2), device)
      result = evaluate(model, test_loader)
      result


{'val_loss': 1.289367914199829, 'val_acc': 0.7705078125}
```

Just as a sanity check, let's verify that this model has the same loss and accuracy on the test set as before.


```
✓ 5s evaluate(model2, test_loader)

{'val_loss': 1.289367914199829, 'val_acc': 0.7705078125}
```

Accuracy vs. No. of epochs




Loss vs. No. of epochs



✓
0s

```
[122] # -----  
# Graded Cell  
# -----  
img, label = test_dataset[1000]  
plt.imshow(img.permute(1, 2, 0))  
print('Label:', training_data.classes[label], ', Predicted:', predict_image(img, model))
```

Label: automobile , Predicted: automobile




↑ ↓ ↶ ↷ ⚙️ 🗑️ ⋮

✓
1s



```
# -----  
# Graded Cell  
# -----  
img, label = test_dataset[2000]  
plt.imshow(img.permute(1, 2, 0))  
print('Label:', training_data.classes[label], ', Predicted:', predict_image(img, model))
```


Label: bird , Predicted: bird




```
✓ 3m [132] history_with_early_stopping = fit_with_early_stopping(num_epochs, lr, model3, train_dl, val_dl, opt_fu

Epoch [0], train_loss: 1.7655, val_loss: 1.4156, val_acc: 0.4529
Epoch [1], train_loss: 1.2532, val_loss: 1.1793, val_acc: 0.5792
Epoch [2], train_loss: 1.0130, val_loss: 0.9405, val_acc: 0.6706
Epoch [3], train_loss: 0.8440, val_loss: 0.8431, val_acc: 0.6999
Epoch [4], train_loss: 0.7211, val_loss: 0.7579, val_acc: 0.7359
Epoch [5], train_loss: 0.6195, val_loss: 0.7542, val_acc: 0.7445
Epoch [6], train_loss: 0.5343, val_loss: 0.7200, val_acc: 0.7463
Epoch [7], train_loss: 0.4599, val_loss: 0.7158, val_acc: 0.7665
Epoch [8], train_loss: 0.3841, val_loss: 0.7487, val_acc: 0.7714
Epoch [9], train_loss: 0.3239, val_loss: 0.7143, val_acc: 0.7722
Epoch [10], train_loss: 0.2595, val_loss: 0.7848, val_acc: 0.7721
Epoch [11], train_loss: 0.2172, val_loss: 0.8628, val_acc: 0.7737
Epoch [12], train_loss: 0.1725, val_loss: 0.9444, val_acc: 0.7636
Epoch [13], train_loss: 0.1548, val_loss: 1.0740, val_acc: 0.7580
Epoch [14], train_loss: 0.1346, val_loss: 1.0676, val_acc: 0.7638
Early stopping at epoch 14
```

We can also plot the validation set accuracies to study how the model improved before it was halted due to the early stopping meta-algorithm. Moreover, we can plot the training and validation losses to study how the effect of overfitting was avoided when we introduced early stopping.





```
[135] # -----  
# Graded Cell  
# -----  
evaluate(model3, test_loader)  
  
{'val_loss': 1.0728938579559326, 'val_acc': 0.764453113079071}
```